

Understanding Clouds from Satellite Images



Sprint 1 Goals

— — —

- Decide which image processing API we are going to use.
- Formulate a goal accuracy to aim for
- Create a working image analyzer for cloud types/density.
- Combine the API and our image analyzing program into one master program.

Steps Taken

We began to look at other Kaggle competitions to get a sense of where to begin. They were also instrumental in understanding how machine learning code works.

<https://www.kaggle.com/c/titanic>

The Kaggle Titanic competition was a great place to start to learn the basic concepts and strategies to solving machine learning problems.

Sprint 1 Outcomes

- Using Keras with Tensorflow backend
 - Why?
 - Higher level API: Very easy to use
 - Supports Multiple backends: Tensorflow, Theano, CNTK
 - Used heavily
- Score to aim for : 0.50 - 0.60
 - Why?
 - Highest Score in competition currently is : 0.677
 - Score is evaluated as: Dice Coefficient

$$\frac{2 * |X \cap Y|}{|X| + |Y|}$$

X is the predicted set of pixels and Y is the ground truth

Sprint 2

— — —

- Results
- Demo
- System Architecture
- Major Components

Results

Results

- Dice coefficient score: 0.577 (*)
- Highest is 0.677

* We did not right all of the code. A **lot** of the code is common knowledge passed down between kaggle kernels.

Demo / Kaggle Notebook

Demo / Kaggle Notebook


— — —

- Link to notebook:

<https://www.kaggle.com/arorashu/cloud-det-nb/data?scriptVersionId=22995066>

High level System Architecture



1. Preprocess Data
2. Define neural net architecture
3. Perform Training : Model Fitting



1. Preprocess Data

How the Input Data looks like

	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q																			
1	Image_Label	EncodedPixels																																		
2	0011165.jpg_Fish	264918	937	266318	937	267718	937	269118	937	270518	937	271918	937	273318	937	274718	937	276118	937	277518	937	278918	937	280318	937	281718	937	283118	937	284518						
3	0011165.jpg_Flower	1355565	1002	1356965	1002	1358365	1002	1359765	1002	1361165	1002	1362565	1002	1363965	1002	1365365	1002	1366765	1002	1368165	1002	1369565	1002	1370965	1002	1372365	1002	1373765	1002	1375165						
4	0011165.jpg_Gravel																																			
5	0011165.jpg_Sugar																																			
6	002be4f.jpg_Fish	233813	878	235213	878	236613	878	238010	881	239410	881	240810	881	242210	881	243610	881	245010	881	246410	881	247810	881	249210	881	250610	881	252010	881	253410						
7	002be4f.jpg_Flower	1339279	519	1340679	519	1342079	519	1343479	519	1344879	519	1346279	519	1347679	519	1349079	519	1350479	519	1351879	519	1353279	519	1354679	519	1356079	519	1357479	519	1358879						
8	002be4f.jpg_Gravel																																			
9	002be4f.jpg_Sugar	67495	350	68895	350	70295	350	71695	350	73095	350	74495	350	75895	350	77295	350	78695	350	80095	350	81495	350	82895	350	84295	350	85695	350	87095	350	88495	350	89895		
10	0031ae9.jpg_Fish	3510	690	4910	690	6310	690	7710	690	9110	690	10510	690	11910	690	13310	690	14710	690	16110	690	17510	690	18910	690	20310	690	21710	690	23110	690	24510	690	25910		
11	0031ae9.jpg_Flower	2047	703	3447	703	4847	703	6247	703	7647	703	9047	703	10447	703	11847	703	13247	703	14647	703	16047	703	17447	703	18847	703	20247	703	21647	703	23047	703	24447	703	25847
12	0031ae9.jpg_Gravel																																			
13	0031ae9.jpg_Sugar	658170	388	659570	388	660970	388	662370	388	663770	388	665170	388	666570	388	667970	388	669370	388	670770	388	672170	388	673570	388	674970	388	676370	388	677770	388	679170	388	680570		



How the Input Data looks like

Training Images: 5546

3.46 GB

Test Images: 3698

2.3 GB

Labels: $5546 * 4 = 22184$

200MB


Time for 10 epochs: around 2 hours - on a GPU!

2. Define neural net architecture

ResNet

— — —

1. Utilizes skip connection
2. Most useful trick to get more performance out of deeper layers
3. We use the Standard ResNet-34



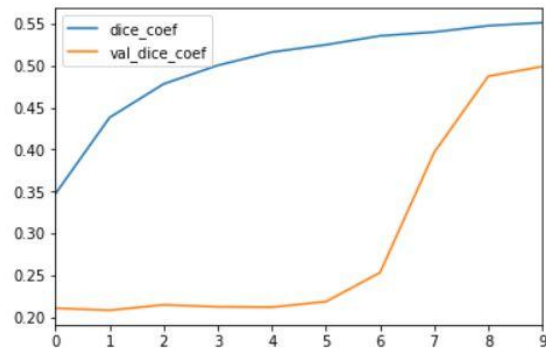

Why ResNet?

1. ResNet are the state of the art in Image Segmentation
2. Libraries like Tensorflow have very efficient implementations
3. ResNets, UNets and their variants are the most in this Kaggle competition

3. Perform Training : Model Fitting

Perform Training : Model Fitting

```
Epoch 1/10  
138/138 [=====] - 800s 6s/step - loss: 1.1424 - dice_coef: 0.3473 - val_loss: 1.9020 - val_dice_coef: 0.2107  
Epoch 2/10  
138/138 [=====] - 814s 6s/step - loss: 0.8962 - dice_coef: 0.4381 - val_loss: 1.8795 - val_dice_coef: 0.2082  
Epoch 3/10  
138/138 [=====] - 763s 6s/step - loss: 0.8361 - dice_coef: 0.4782 - val_loss: 1.6433 - val_dice_coef: 0.2147  
Epoch 4/10  
138/138 [=====] - 826s 6s/step - loss: 0.8104 - dice_coef: 0.5003 - val_loss: 2.2932 - val_dice_coef: 0.2123  
Epoch 5/10  
138/138 [=====] - 754s 5s/step - loss: 0.7897 - dice_coef: 0.5161 - val_loss: 2.0335 - val_dice_coef: 0.2119  
Epoch 6/10  
138/138 [=====] - 819s 6s/step - loss: 0.7795 - dice_coef: 0.5249 - val_loss: 2.2183 - val_dice_coef: 0.2185  
Epoch 7/10  
138/138 [=====] - 764s 6s/step - loss: 0.7658 - dice_coef: 0.5355 - val_loss: 1.9942 - val_dice_coef: 0.2532  
Epoch 8/10  
138/138 [=====] - 755s 5s/step - loss: 0.7597 - dice_coef: 0.5400 - val_loss: 1.1998 - val_dice_coef: 0.3965  
Epoch 9/10  
138/138 [=====] - 764s 6s/step - loss: 0.7486 - dice_coef: 0.5476 - val_loss: 0.8975 - val_dice_coef: 0.4873  
Epoch 10/10  
138/138 [=====] - 739s 5s/step - loss: 0.7443 - dice_coef: 0.5512 - val_loss: 0.9633 - val_dice_coef: 0.4991
```



+ Code

+ Markdown